

# Master's Thesis

## Neurosymbolic AI for Social Cognition

Jarne Demoen

Vrije Universiteit Brussel

18/10/2025

## 1 Goal and Overview

We build a **Mixture-of-Experts (MoE)** model combining four emotion predictors:

- Three **face experts**:  $e \in \{0, 1, 2\}$ , each outputting 7 emotion logits  $\mathbf{z}^{(e)} \in \mathbb{R}^7$ ,
- One **scene expert**:  $\mathbf{z}^{(\text{scn})} \in \mathbb{R}^7$ .

A learned **gating network** computes per-image weights  $\mathbf{w}(\mathbf{x}) \in \Delta^3$  (the 4-dimensional simplex), and the experts are mixed on the probability level:

$$\mathbf{p}_{\text{neural}}(\mathbf{x}) = \sum_{e \in \{0, 1, 2, \text{scn}\}} w_e(\mathbf{x}) \underbrace{\text{softmax}(\mathbf{z}^{(e)}(\mathbf{x}))}_{\mathbf{p}^{(e)}(\mathbf{x})}.$$

This neural mixture is then combined symbolically with domain **prior knowledge** inside DeepProbLog:

$$p(E \mid \mathbf{x}, \text{Meta}) \propto p_{\text{neural}}(E \mid \mathbf{x}) \cdot p_{\text{prior}}(E \mid \text{Meta}),$$

which is exactly a *product-of-experts*. The final predicate `final_emo/4` corresponds to this product and is the one used for training and inference.

## 2 Notation and Shapes

For a mini-batch:

$B$  = batch size,  $C = 7$  emotion classes,  $F = 3$  face slots.

$\mathbf{z}_{\text{face}} \in \mathbb{R}^{B \times F \times C}$ ,  $\mathbf{z}_{\text{scn}} \in \mathbb{R}^{B \times C}$ .

$\mathbf{p}_{\text{face}} = \text{softmax}(\mathbf{z}_{\text{face}}, -1) \in \mathbb{R}^{B \times F \times C}$ ,  $\mathbf{p}_{\text{scn}} = \text{softmax}(\mathbf{z}_{\text{scn}}, -1) \in \mathbb{R}^{B \times C}$ .

Gating weights:  $\mathbf{w} \in \mathbb{R}^{B \times 4}$  with  $\sum_e w_{b,e} = 1$  and  $w_{b,e} \geq 0$ .

## 3 Gating Inputs (Features)

The gate should learn which expert is reliable for each image. We feed it compact indicators of confidence and context.

## Per expert (face or scene)

$$m^{(e)} = \max_c p_c^{(e)} \quad (\text{maximum probability, confidence})$$

$$h^{(e)} = - \sum_c p_c^{(e)} \log p_c^{(e)} \quad (\text{entropy, uncertainty})$$

## Additional contextual signals

- `face_present[e] ∈ {0, 1}` – whether face slot  $e$  exists,
- `num_faces ∈ {0, 1, 2, 3}`,
- Optionally, detection confidence, blur level, or similar.

The concatenated gate input per image may thus include:

$$\text{gate\_in} = [\mathbf{z}_{\text{face}} (\text{flattened}), \mathbf{z}_{\text{scn}}, \mathbf{m}, \mathbf{h}, \text{face\_present}, \text{num\_faces}],$$

but you can begin with only  $(\mathbf{m}, \mathbf{h}, \text{flags})$  for simplicity.

## 4 Gating Network and Probability Mixture

The gate is a simple MLP producing 4 weights:

$$\mathbf{w}(\mathbf{x}) = \text{softmax}(f_\theta(\text{gate\_features}(\mathbf{x}))) \in \mathbb{R}^4.$$

The neural mixture:

$$\mathbf{p}_{\text{neural}}(\mathbf{x}) = \sum_{e \in \{0,1,2,\text{scn}\}} w_e(\mathbf{x}) \mathbf{p}^{(e)}(\mathbf{x}), \quad \sum_e w_e(\mathbf{x}) = 1.$$

*# probabilities per expert*

```
p_face = torch.softmax(z_face, dim=-1) # (B, 3, 7)
p_scene = torch.softmax(z_scene, dim=-1) # (B, 7)
```

*# gate features*

```
max_face = p_face.max(dim=-1).values
ent_face = -(p_face * (p_face.clamp_min(1e-12)).log()).sum(dim=-1)
max_scene = p_scene.max(dim=-1).values
ent_scene = -(p_scene * (p_scene.clamp_min(1e-12)).log()).sum(dim=-1)
```

```
feat_list = [max_face, ent_face,
              max_scene.unsqueeze(-1), ent_scene.unsqueeze(-1),
              face_present.float(), num_faces]
gate_in = torch.cat([t.reshape(t.size(0), -1) for t in feat_list], dim=1)
# (B,D)
```

*# MLP -> weights*

```
w = torch.softmax(gate_mlp(gate_in), dim=-1) # (B,4)
```

```
# probability mixture
p_mix = (
    (w[:,0].unsqueeze(-1).unsqueeze(-1) * p_face[:,0]) +
    (w[:,1].unsqueeze(-1).unsqueeze(-1) * p_face[:,1]) +
    (w[:,2].unsqueeze(-1).unsqueeze(-1) * p_face[:,2]) +
    (w[:,3].unsqueeze(-1) * p_scene)
)
```

## 5 Symbolic Priors in DeepProbLog

**Implementation sketch (PyTorch).** Symbolic priors encode domain knowledge  $p_{\text{prior}}(E \mid \text{Meta})$ . In ProbLog, conjunction means multiplication of probabilities, so combining a neural predicate with a prior predicate yields a product-of-experts model.

**Formally:**

$$p(E \mid \mathbf{x}, \text{Meta}) \propto p_{\text{neural}}(E \mid \mathbf{x}) \cdot p_{\text{prior}}(E \mid \text{Meta}).$$

```
% Neural mixture (linked to your PyTorch MoE)
nn(moe_net, [Faces, Scene], E, [0,1,2,3,4,5,6]) :: neural_emo(Faces, Scene,

% Metadata facts (per example):
% num_faces(MetaId, N).    scene_tag(MetaId, Tag).
% Example:
%   num_faces(meta123, 0).
%   scene_tag(meta123, party).

% Soft baseline prior (every class gets some mass)
0.20::prior_emo(_, 0). 0.20::prior_emo(_, 1). 0.20::prior_emo(_, 2).
0.20::prior_emo(_, 3). 0.20::prior_emo(_, 4). 0.20::prior_emo(_, 5).
0.20::prior_emo(_, 6).

% Rules:
0.10::prior_emo(M, 0) :- num_faces(M, 0). % damp anger if no faces
0.10::prior_emo(M, 1) :- num_faces(M, 0). % damp disgust
0.10::prior_emo(M, 5) :- num_faces(M, 0). % damp surprise

0.60::prior_emo(M, 3) :- scene_tag(M, party). % boost happy
0.15::prior_emo(M, 4) :- scene_tag(M, party). % damp sad

% Combine neural and prior via conjunction (product of probs):
final_emo(Faces, Scene, Meta, E) :-
    neural_emo(Faces, Scene, E),
    prior_emo(Meta, E).
```

**Training:** Train and query `final_emo/4`, not `neural_emo/3`.

## 6 Training Objective

DeepProbLog optimizes the negative log-likelihood of the final query:

$$\mathcal{L} = -\log p(E^* \mid \mathbf{x}, \text{Meta}),$$

where  $p$  is the probability returned by the probabilistic reasoning engine for `final_emo`.

### Implementation steps.

1. Build the PyTorch MoE (experts  $\rightarrow$  logits, gate  $\rightarrow \mathbf{w}$ , mix  $\rightarrow \mathbf{p}_{\text{neural}}$ ).
2. Link `moe_net` to DeepProbLog via `nn(...)` predicate.
3. Add metadata facts: `num_faces/2`, `scene_tag/2`, etc.
4. Define symbolic `prior_emo/2`.
5. Train using `final_emo(Faces, Scene, Meta, E_gold)` queries.

## 7 Worked Example

Suppose (for one image):

- $\mathbf{p}^{(0)} = [.30, .05, .05, .30, .10, .10, .10]$ ,
- $\mathbf{p}^{(1)} = [.25, .05, .05, .35, .10, .10, .10]$ ,
- $\mathbf{p}^{(2)} = [.10, .10, .10, .50, .05, .10, .05]$ ,
- $\mathbf{p}^{(\text{scn})} = [.05, .05, .10, .45, .15, .10, .10]$ ,
- $\mathbf{w} = [0.2, 0.2, 0.2, 0.4]$ .

Then

$$\mathbf{p}_{\text{neural}} = 0.2\mathbf{p}^{(0)} + 0.2\mathbf{p}^{(1)} + 0.2\mathbf{p}^{(2)} + 0.4\mathbf{p}^{(\text{scn})}.$$

If metadata indicates `num_faces=0` and `scene_tag=party`, the symbolic priors amplify `happy` and damp `sad/anger/disgust/surprise`. The final DeepProbLog inference multiplies and renormalizes, shifting probability mass towards `happy` in an explainable way.

## 8 Practical Stability Tips

- Clamp in logs: use `clamp_min(1e-12)` for numerical stability.
- Normalize gating inputs: entropies  $\in [0, \log C]$ .
- If a face slot is empty: `face_present=0`, feed a flat or neutral distribution.
- Use small gate MLP (1–2 layers, 64 hidden) with weight decay to avoid overfitting.

## 9 Interface Summary

- **Neural side:** `nn(moe_net, [Faces, Scene], E, [0..6])` returns  $\mathbf{p}_{\text{neural}}$ .
- **Data:**
  - Faces: tensor source with 3 face crops or features.
  - Scene: scene image tensor.
  - Meta: symbolic facts per example.

## 10 Implement Now vs Later

### Implement Now (core system)

1. Experts  $\rightarrow$  logits  $\rightarrow$  softmax  $\rightarrow$  per-expert  $\mathbf{p}^{(e)}$ .
2. Gate MLP with **max-prob**, **entropy**, **face\_present**, **num\_faces**.
3. Probability mixture  $\mathbf{p}_{\text{neural}} = \sum_e w_e \mathbf{p}^{(e)}$ .
4. Symbolic priors (`prior_emo`) and training on `final_emo`.

### Add Later (extensions)

- **Temperature calibration** per expert: learn  $\tau_e$  so that  $\mathbf{p}^{(e)} = \text{softmax}(\mathbf{z}^{(e)}/\tau_e)$ .
- **Agreement features:** Jensen-Shannon or KL divergence between experts.
- **Logit mixture (log-sum-exp):**

$$\mathbf{z}_{\text{mix}} = \log\left(\sum_e w_e e^{\mathbf{z}^{(e)}}\right), \quad \mathbf{p}_{\text{neural}} = \text{softmax}(\mathbf{z}_{\text{mix}}).$$

- Richer priors (object tags, time, context) and symbolic features for the gate.

## 11 Checklist

1. Gate outputs 4 weights per sample (softmax-normalized).
2. Shape consistency:  $(B, 3, 7), (B, 7) \rightarrow (B, 7)$ .
3. Priors are soft (no hard zeros).
4. Train and query `final_emo/4`.
5. Log which symbolic rules fired for explainability.